

HDFS Architecture

1.

HDFS is strictly POSIX compliant.



True



False

2.

Following issues may be caused by lot of small files in HDFS



NameNode memory usage increases significantly



Network load decreases



Number of map tasks need to process the same amount of data will be larger.



I/O rate will be faster

3.

You are writing a 10GB file to a HDFS filesystem with a default block size of 128MB. How many blocks will the file be broken into?



1280



160



80



480

4.

You are writing a 10GB file into HDFS with a replication of 2 and block size of 64MB. How much total disk space will this file use?



20GB



128GB



10GB



30GB

5.

What is the first step in a write process from a HDFS client?



Immediately contact the NameNode



Start locally caching the data that needs to be written and then contact NameNode



Contact the nearest DataNode



Contact all the DataNodes involved in replication

6.

HDFS NameNode is not rack aware when it places the replica blocks.



True



False